

# ECE 408- Applied Parallel Programming Project

Hwanseo Choi (hwanseo2),  
Simeng Liu (simengl2),  
Chandan Vempati (vempati2).

April 17, 2018

# Chapter 1

## Milestone I

### Question 1

TEAM INFO:

**Team name :** o \_ \_ - O \_ - \_ O \_ \_ - O ,

**Team Members:**

Hwanseo Choi (hwanseo2),

Simeng Liu (simengl2),

Chandan Vempati(vempati2),

.

### Question 2

The list of kernels and some their corresponding parameters which were measured are shown in Table 1.

Name	Time(%)	Time	Calls
void fermiPlusCgemmLDS128 _batched	34.09	118.48 ms	9
void cudnn::detail::implicit _convolve _sgemm	27.02	93.917 ms	1
void fft2d _c2r _32x32	12.62	43.847 ms	9
sgemm _sm35 _ldg _tn _128x8x256x16x32	8.20	28.512 ms	1
CUDA _memcpy _HtoD	6.46	22.445 ms	14
void cudnn::detail::activation_fw _4d _kernel	4.07	14.158 ms	2
void cudnn::detail::pooling_fw _4d _kernel	3.82	13.292 ms	1
void fft2d _r2c _32x32	1.72	5.9651 ms	9
sgemm _sm35 _ldg _tn _64x16x128x8x32	1.17	4.0583 ms	1
void mshadow::cuda::MapPlanLargeKernel	0.37	1.2844 ms	1
void mshadow::cuda::SoftmaxKernel	0.32	1.1046 ms	1
void mshadow::cuda::MapPlanKernel	0.05	177.02 $\mu$ s	13
void mshadow::cuda::MapPlanKernel	0.04	146.34 $\mu$ s	2
sgemm _sm35 _ldg _tn _32x16x64x8x16	0.04	130.11 $\mu$ s	1
void mshadow::cuda::MapPlanKernel	0.01	22.399 $\mu$ s	1
void fft2d _r2c _32x32	0.01	20.671 $\mu$ s	1
CUDA _memcpy _DtoH	0.00	9.9200 $\mu$ s	1

Table 1.1: CUDA kernel calls and their corresponding parameters

### Question 3

The list of APIs and some their corresponding parameters which were measured are shown in Table 1.

Name	Time(%)	Time	Calls
cudaStreamCreateWithFlags	43.55	1.92546 s	18
cudaFree	27.11	1.19848 s	10
cudaMemGetInfo	20.70	915.17 ms	27
cudaStreamSynchronize	7.31	323.39 ms	29
cudaMemcpy2DAsync	1.01	44.605 ms	9
cudaMalloc	0.16	7.2049 ms	45
cudaStreamCreate	0.03	1.5196 ms	4
cuDeviceTotalMem	0.03	1.3522 ms	4
cuDeviceGetAttribute	0.03	1.1863 ms	352
cudaEventCreateWithFlags	0.02	1.0891 ms	114
cudaLaunch	0.02	728.51 $\mu$ s	53
cudaMemcpy	0.01	405.96 $\mu$ s	6
cudaSetupArgument	0.01	352.84 $\mu$ s	619
cudaDeviceGetAttribute	0.00	135.99 $\mu$ s	116
cuDeviceGetName	0.00	102.82 $\mu$ s	4
cudaSetDevice	0.00	82.812 $\mu$ s	35
cudaStreamWaitEvent	0.00	55.338 $\mu$ s	27
cudaStreamCreateWithPriority	0.00	50.246 $\mu$ s	2
cudaConfigureCall	0.00	48.487 $\mu$ s	53
cudaGetDevice	0.00	26.507 $\mu$ s	10
cudaEventRecord	0.00	21.586 $\mu$ s	12
cudaGetLastError	0.00	20.909 $\mu$ s	34
cudaBindTexture	0.00	15.628 $\mu$ s	1
cudaPeekAtLastError	0.00	12.229 $\mu$ s	18
cuDeviceGetCount	0.00	6.8830 $\mu$ s	6
cudaEventCreate	0.00	5.9220 $\mu$ s	1
cudaStreamGetPriority	0.00	5.8400 $\mu$ s	1
cuDeviceGet	0.00	5.0390 $\mu$ s	6
cudaDeviceGetStreamPriorityRange	0.00	4.9430 $\mu$ s	2
cuInit	0.00	3.8850 $\mu$ s	3
cuDriverGetVersion	0.00	2.6300 $\mu$ s	3
cudaEventDestroy	0.00	2.4120 $\mu$ s	1
cudaGetDeviceCount	0.00	2.1210 $\mu$ s	1
cudaUnbindTexture	0.00	1.8500 $\mu$ s	1

Table 1.2: CUDA APIs calls and their corresponding parameters

## Question 4

**Application Program Interfaces** or **APIs** is a set of subroutine definitions that is provided by NVIDIA as part of the CUDA tool and is in-charge of connecting host and device for various purposes. For example, copy memory from host to device or the other direction.

**Kernels** are the custom code that is defined by the user and are executed  $N$  times in parallel by  $N$  different CUDA threads, as opposed to only once like regular C functions.

## Question 5

**Output of rai running MXNET on the CPU is as follows:**

```
Loading fashion-mnist data...
done
Loading model...
done
New Inference
EvalMetric:  'accuracy':  0.8444
```

## Question 6

The program run time on the CPU is **12.74 s**.

## Question 7

**Output of rai running MXNET on the GPU is as follows:**

```
Running /usr/bin/time python m1.2.py
Loading fashion-mnist data...
done
Loading model...
src/operator/././ cudnn _algoreg-inl.h:112: Running performance tests to
find the best convolution algorithm, this can take a while...(setting env
variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done
New Inference
EvalMetric:  'accuracy':  0.8444
```

## Question 8

The program run time on the GPU is **2.13 s**.

# Chapter 2

## Milestone II

### Question 1

The list of total execution time for all parameters is shown in Table 3.

Number of Images	Time (s)
10000 (default)	30.58 user 1.55 system 0:30.03 elapsed
1000	1.06 user 0.61 system 0:01.02 elapsed
100	0.70 user 0.48 system 0:00.72 elapsed

Table 2.1: Total execution times and their corresponding parameters

### Question 2

The list of op time for all parameters is shown in Table 4.

Number of Images	Op Time 1 (s)	Op Time 2 (s)
10000 (default)	6.61	19.48
1000	0.07	0.20
100	0.01	0.02

Table 2.2: Total execution times and their corresponding parameters

The Op Time scales linearly with the number of images.

# Chapter 3

## Milestone III

The output of the parallelized code is as follows:

```
* Running python m3.1.py
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 0.293913
Op Time: 0.720379
Correctness: 0.8451 Model: ece408
* Running python m3.1.py 10
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 0.000235
Op Time: 0.000703
Correctness: 1.0 Model: ece408
* Running python m3.1.py 100
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 0.002745
Op Time: 0.008985
Correctness: 0.88 Model: ece408
```

```

void mshadow::cuda::MapPlanKernel, float>,
mshadow::expr::Plan, mshadow::Tensor, float, int=1>,
float>>(mshadow::gpu, unsigned int, mshadow::Shape, int=4)
1.25\% 198.01us 14 14.143us 2.0480us 75.071us void
mshadow::cuda::MapPlanKernel, float>, mshadow::expr::Plan,
float>>(mshadow::gpu, unsigned int, mshadow::Shape, int=2)
1.01% 160.13us 2 80.063us 8.2560us 151.87us void
cudnn::detail::activation_fw_4d_kernel>(cudnnTensorStruct, float
const , cudnn::detail::activation_fw_4d_kernel>,
cudnnTensorStruct, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
0.89% 141.66us 1 141.66us 141.66us 141.66us void
cudnn::detail::pooling_fw_4d_kernel, int=0>\\(cudnnTensorStruct, float const
, cudnn::detail::pooling_fw_4d_kernel,
int=0>, cudnnTensorStruct, cudnnPoolingStruct, float, cudnnPoolingStruct,
int, cudnn::reduced_divisor, float)
0.25% 40.255us 1 40.255us 40.255us 40.255us sgemm_sm35_ldg_tn_32x16x64x8x16
0.12% 18.560us 1 18.560us 18.560us 18.560us void mshadow::cuda::SoftmaxKernel,
float>, mshadow::expr::Plan, float>>(mshadow::gpu,
int=2, unsigned int)
0.08% 12.000us 1 12.000us 12.000us 12.000us void
mshadow::cuda::MapPlanKernel, float>,
mshadow::expr::Plan, float, int=3, bool=1, int=2>,
float
(mshadow::gpu, unsigned int, mshadow::Shape, int=2)
0.05% 7.3590us 2 3.6790us 3.0080us 4.3510us void
mshadow::cuda::MapPlanKernel, float>,
mshadow::expr::Plan, float, int=2, int=1>, float>>(mshadow::gpu, unsigned
int, mshadow::Shape, int=2)
0.04% 5.8870us 1 5.8870us 5.8870us 5.8870us [CUDA memcpy DtoH]
0.02% 3.8720us 1 3.8720us 3.8720us 3.8720us void
scal_kernel
(cublasTransposeParams, float const , float, float const *)
0.01% 1.5040us 1 1.5040us 1.5040us 1.5040us [CUDA memset]

```

The subsequent output has been tabulated. We refer you to Table (3.1).



Name	Time (s)	Calls	Avg (ms)	Min ( $\mu$ s)	Max (ms)
cudaStreamCreateWithFlags	1.54822s	16	96.764	19.482	773.80
cudaFree	1.18079s	10	118.08	1.1940	315.38
cudaMemGetInfo	1.00035s	27	37.050	337.28	991.24
cudaDeviceSynchronize	13.589	6	2.2648	7.9570	10.192
cudaMalloc	3.6162	45	80.359	12.882	333.55
cudaMemcpy2DAsync	2.1451	9	238.35	14.825	842.96
cudaStreamSynchronize	1.7335	29	59.775	6.4650	623.37
cuDeviceTotalMem	1.4029	4	350.73	340.77	375.89
cuDeviceGetAttribute	973.49	352	2.7650	516e-3	72.454
cudaEventCreateWithFlags	796.00	112	7.1070	900e-3	308.82
cudaLaunch	580.81	28	20.743	8.5320	55.081
cudaMemcpy	420.03	6	70.005	26.248	151.75
cudaStreamCreate	207.98	4	51.996	24.159	81.876
cuDeviceGetName	119.21	4	29.801	19.015	38.011
cudaSetupArgument	107.69	158	681e-3	523e-3	1.5900
cudaDeviceGetAttribute	99.049	104	952e-3	694e-3	2.3270
cudaSetDevice	91.029	34	2.6770	927e-3	9.8460
cudaMemsetAsync	50.096	1	50.096	50.096	50.096
cudaStreamCreateWithPriority	40.683	2	20.341	20.032	20.651
cudaConfigureCall	39.980	28	1.4270	672e-3	3.5080
cudaGetDevice	33.473	10	3.3470	1.5450	8.5850
cudaPeekAtLastError	16.071	20	803e-3	608e-3	1.0910
cudaGetLastError	8.0830	8	1.0100	557e-3	2.6470
cudaEventQuery	7.4510	1	7.4510	7.4510	7.4510
cuDeviceGetCount	5.9890	6	998e-3	517e-3	2.1960
cuDeviceGet	5.8330	6	972e-3	651e-3	1.2940
cudaDeviceGetStreamPriorityRange	3.9960	2	1.9980	1.5990	2.3970
cuInit	3.8320	3	1.2770	1.1680	1.4220
cuDriverGetVersion	3.1750	3	1.0580	992e-3	1.1300
cudaEventRecord	2.6950	1	2.6950	2.6950	2.6950
cudaGetDeviceCount	2.1030	1	2.1030	2.1030	2.1030

Table 3.1: CUDA APIs calls and their corresponding parameters

**The parameters corresponding to the kernel call are as follows:**

Time(%) Time Calls Avg Min Max Name

92.46% 1.32146s 2 660.73ms 307.68ms 1.01378s mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

★ Running nvprof python m3.1.py 10

48.77% 1.3694ms 2 684.69us 319.74us 1.0496ms mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

★ Running nvprof python m3.1.py 100

83.74% 13.262ms 2 6.6312ms 3.0892ms 10.173ms mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

6.70% 1.0615ms 14 75.818us 1.5680us 813.94us [CUDA memcpy HtoD]

★ Running nvprof python m3.1.py

Time(%) Time Calls Avg Min Max Name

92.46% 1.32146s 2 660.73ms 307.68ms 1.01378s mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

★ Running nvprof python m3.1.py 10

48.77% 1.3694ms 2 684.69us 319.74us 1.0496ms mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

★ Running nvprof python m3.1.py 100

83.74% 13.262ms 2 6.6312ms 3.0892ms 10.173ms mxnet::op::forward\_kernel(float\*, float const \*, float const \*, int, int, int, int, int, int)

6.70% 1.0615ms 14 75.818us 1.5680us 813.94us [CUDA memcpy HtoD]